

# ecg\_rhythm\_classification

December 12, 2024

## 1 ECG Rhythm Classification

The goal of this exercise is to train a neural network model to classify different cardiac rhythm from single-lead ECG signals. The ECG signals we will use are a subset of the large scale 12-lead electrocardiogram database for arrhythmia study (<https://physionet.org/content/ecg-arrhythmia/1.0.0/>).

The subset includes the following cardiac rhythms:

- Atrial fibrillation
- Atrial flutter
- Normal sinus rhythm
- Sinus bradycardia
- Sinus tachycardia

There are 1500 single-lead ECG signals (lead II) for each rhythm.

First, we import all required packages, define global constants, and seed the random number generators to obtain reproducible results.

```
[1]: %matplotlib widget

import collections
import itertools
import logging
import operator
import pathlib
import warnings

import IPython.display
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pytorch_lightning as pl
import sklearn.metrics
import sklearn.model_selection
import sklearn.preprocessing
import torch
```

```

DATA_FILE = pathlib.Path('../data/ecg_rhythms.npz')
LOG_DIRECTORY = pathlib.Path('../logs/ecg_rhythm_classification')

# Disable logging for PyTorch Lightning to avoid too long outputs.
logging.getLogger('pytorch_lightning').setLevel(logging.ERROR)

# Seed random number generators for reproducible results.
pl.seed_everything(42)

```

Seed set to 42

[1]: 42

Then, we load the ECG signals and the corresponding rhythm annotations.

```

[2]: def load_data():
    with np.load(DATA_FILE) as data:
        signals = data['signals']
        rhythms = data['rhythms']
        fs = data['fs'].item()
    return signals, rhythms, fs

signals, rhythms, fs = load_data()

IPython.display.display(pd.DataFrame(sorted(collections.Counter(rhythms).
    ↪items()), columns=['rhythm', 'count']))

```

	rhythm	count
0	atrial_fibrillation	1500
1	atrial_flutter	1500
2	normal_sinus_rhythm	1500
3	sinus_bradycardia	1500
4	sinus_tachycardia	1500

Here are a few examples of ECG signals.

```


[3]: def plot_ecg_examples(signals, rhythms, fs):
    time = np.arange(signals.shape[-1]) / fs
    labels = np.unique(rhythms)
    indices = [np.random.choice(np.flatnonzero(rhythms == label)) for label in_
    ↪labels]

    fig, axes = plt.subplots(len(labels), 1, sharex='all',_
    ↪constrained_layout=True)
    for ax, index in zip(axes.flat, indices):
        label = rhythms[index].replace('_', ' ').capitalize()
        plt.sca(ax)

```

```
plt.plot(time, signals[index].T, linewidth=1, label=label)
plt.grid()
plt.ylabel('ECG')
plt.legend(loc='upper right')
plt.xlabel('Time [s]')
```

```
plot_ecg_examples(signals, rhythms, fs)
```



### Question 1

Visually, what are the differences between the different rhythms?

#### Answer

*Atrial fibrillation*, as it has been shown many times in the other lab sessions, shows a lot of irregularities in its PR-intervals and a shift of R and S points : instead of being respectively around 0.5 and -0.1, their values are around 0.25 and -0.25. There is also an increase in heartbeat frequency.

The *atrial flutter* signal shows a huge increase in heartbeat frequency (much bigger than the one in atrial fibrillation, nearly twice the one of the normal signal).

*Sinus bradycardia* is the opposite and shows a decrease in frequency. There is also an increase in

amplitude of the R-peaks.

*Sinus tachycardia* has a signal frequency slightly lower than the *atrial flutter* one, with one every four PQRS complexes showing really extreme behaviours (extreme amplitudes of R and S points).

Then, we split that data into subsets for training, validation, and testing stratified by rhythms.

```
[4]: def split_data(rhythms):
    n = rhythms.size
    splitter = sklearn.model_selection.StratifiedKFold(n_splits=5)
    indices = list(map(operator.itemgetter(1), splitter.split(np.zeros((n, 1)),
    ↪rhythms)))
    i_train = np.hstack(indices[:-2])
    i_val = indices[-2]
    i_test = indices[-1]

    assert np.intersect1d(i_train, i_val).size == 0
    assert np.intersect1d(i_train, i_test).size == 0
    assert np.intersect1d(i_val, i_test).size == 0
    assert np.all(np.sort(np.hstack((i_train, i_val, i_test))) == np.arange(n))

    return i_train, i_val, i_test
```

```
i_train, i_val, i_test = split_data(rhythms)
```

```
def build_summary(rhythms, indices):
    labels = np.unique(rhythms)
    data = []
    for subset, i in indices:
        y = rhythms[i]
        data.append({'subset': subset, 'total_count': y.size})
        for label in labels:
            data[-1][f'{label}_count'] = np.sum(y == label)
    return pd.DataFrame(data)
```

```
IPython.display.display(build_summary(rhythms, (('train', i_train), ('val',
    ↪i_val), ('test', i_test))))
```

	subset	total_count	atrial_fibrillation_count	atrial_flutter_count	\
0	train	4500	900	900	
1	val	1500	300	300	
2	test	1500	300	300	

	normal_sinus_rhythm_count	sinus_bradycardia_count	sinus_tachycardia_count
0	900	900	900
1	300	300	300

The final preprocessing steps are to scale the ECG signals so that they have approximately unit variance and to encode the rhythm labels with one-hot encoding.

```
[5]: def compute_scaling(signals):
    sigma = np.std(signals)
    return 1.0 / sigma

alpha = compute_scaling(signals[i_train])
signals *= alpha

def encode_rhythms(rhythms):
    categories = [np.unique(rhythms)]
    encoder = sklearn.preprocessing.OneHotEncoder(categories=categories,
    ↪sparse_output=False)
    return encoder.fit_transform(rhythms[:, None])

encoded_rhythms = encode_rhythms(rhythms)

def print_encoded_rhythms(rhythms, encoded_rhythms, n=10):
    df = pd.DataFrame(encoded_rhythms, columns=np.unique(rhythms))
    df.insert(0, 'rhythm', rhythms)
    IPython.display.display(df.head(n))

print_encoded_rhythms(rhythms, encoded_rhythms)
```

	rhythm	atrial_fibrillation	atrial_flutter \
0	atrial_fibrillation	1.0	0.0
1	sinus_bradycardia	0.0	0.0
2	sinus_bradycardia	0.0	0.0
3	atrial_flutter	0.0	1.0
4	sinus_bradycardia	0.0	0.0
5	atrial_fibrillation	1.0	0.0
6	normal_sinus_rhythm	0.0	0.0
7	sinus_bradycardia	0.0	0.0
8	sinus_bradycardia	0.0	0.0
9	sinus_bradycardia	0.0	0.0

	normal_sinus_rhythm	sinus_bradycardia	sinus_tachycardia
0	0.0	0.0	0.0
1	0.0	1.0	0.0
2	0.0	1.0	0.0

3	0.0	0.0	0.0
4	0.0	1.0	0.0
5	0.0	0.0	0.0
6	1.0	0.0	0.0
7	0.0	1.0	0.0
8	0.0	1.0	0.0
9	0.0	1.0	0.0

We define a class and few utility functions for training and evaluating models.

```
[6]: class Classifier(pl.LightningModule):

    def __init__(self, model, learning_rate=0.001):
        super().__init__()
        self.save_hyperparameters(ignore=['model'])
        self.model = model
        self.learning_rate = learning_rate
        self.example_input_array = torch.zeros((1,) + self.model.input_shape)

    def configure_optimizers(self):
        return torch.optim.Adam(self.parameters(), lr=self.learning_rate)

    def forward(self, x):
        return self.model(x)

    def training_step(self, batch, batch_idx):
        return self._run_step(batch, 'train')

    def validation_step(self, batch, batch_idx):
        self._run_step(batch, 'val')

    def test_step(self, batch, batch_idx):
        self._run_step(batch, 'test')

    def predict_step(self, batch, batch_idx, dataloader_idx=0):
        x, y = batch
        return self.model(x)

    def _run_step(self, batch, subset):
        x, y = batch
        logits = self.model(x)
        loss = torch.nn.functional.cross_entropy(logits, y)
        acc = (torch.argmax(y, 1) == torch.argmax(logits, 1)).float().mean()
        self.log_dict({
            f'{subset}_loss': loss,
            f'{subset}_acc': acc,
        }, on_step=False, on_epoch=True, prog_bar=True)
        return loss
```

```

def build_loader(*tensors, batch_size=100, shuffle=False, n_workers=0):
    dataset = torch.utils.data.TensorDataset(*map(torch.Tensor, tensors))
    return torch.utils.data.DataLoader(
        dataset=dataset,
        batch_size=batch_size,
        shuffle=shuffle,
        num_workers=n_workers,
    )

def train_model(name, model, x, y, i_train, i_val, learning_rate=0.001,
    ↪batch_size=100, n_epochs=10):
    train_loader = build_loader(x[i_train], y[i_train], batch_size=batch_size,
    ↪shuffle=True)
    val_loader = build_loader(x[i_val], y[i_val], batch_size=batch_size)
    classifier = Classifier(model, learning_rate)
    print(pl.utilities.model_summary.ModelSummary(classifier, max_depth=-1))

    with warnings.catch_warnings():
        warnings.simplefilter('ignore')
        trainer = pl.Trainer(
            default_root_dir=LOG_DIRECTORY,
            logger=pl.loggers.TensorBoardLogger(LOG_DIRECTORY, name),
            enable_model_summary=False,
            max_epochs=n_epochs,
        )
        trainer.fit(classifier, train_loader, val_loader)

    return classifier

def evaluate_model(model, x, y, i_train, i_val, i_test, batch_size=100):
    loader = build_loader(x, y, batch_size=batch_size)

    with warnings.catch_warnings():
        warnings.simplefilter('ignore')
        trainer = pl.Trainer(
            default_root_dir=LOG_DIRECTORY,
            logger=False,
            enable_progress_bar=False,
            enable_model_summary=False,
        )
        z = trainer.predict(model, loader)
        z = np.vstack([u.numpy() for u in z])

```

```

references = np.argmax(y, axis=1)
predictions = np.argmax(z, axis=1)
matrices = {}
for subset, indices in (('train', i_train), ('val', i_val), ('test', i_test)):
    matrices[subset] = sklearn.metrics.confusion_matrix(
        references[indices],
        predictions[indices],
    )

return matrices

```

We start TensorBoard to visualize performance metrics during training.

If you prefer to view TensorBoard in a separate window, you can open <http://localhost:6006/> in your web browser.

```

[7]: %reload_ext tensorboard
      %tensorboard --logdir ../logs/ecg_rhythm_classification --port 6006

```

<IPython.core.display.HTML object>

We define a convolutional neural network.

```

[8]: class CnnModel(torch.nn.Module):

      def __init__(self, input_shape, output_shape, kernel_size=5):
          super().__init__()
          self.input_shape = input_shape
          self.output_shape = output_shape
          self.layers = torch.nn.Sequential(
              torch.nn.Conv1d(self.input_shape[0], 8, kernel_size,
                  padding='same'),
              torch.nn.BatchNorm1d(8),
              torch.nn.ReLU(),
              torch.nn.MaxPool1d(2),

              torch.nn.Conv1d(8, 16, kernel_size, padding='same'),
              torch.nn.BatchNorm1d(16),
              torch.nn.ReLU(),
              torch.nn.MaxPool1d(2),

              torch.nn.Conv1d(16, 32, kernel_size, padding='same'),
              torch.nn.BatchNorm1d(32),
              torch.nn.ReLU(),
              torch.nn.MaxPool1d(2),

              torch.nn.Conv1d(32, 64, kernel_size, padding='same'),
              torch.nn.BatchNorm1d(64),

```



```

        torch.nn.ReLU(),
        torch.nn.AdaptiveAvgPool1d(1),

        torch.nn.Flatten(),
        torch.nn.Linear(64, self.output_shape[0]),
    )

    def forward(self, x):
        return self.layers(x)

```

Then, we train and evaluate this model.

```

[9]: input_shape = signals.shape[1:]
    output_shape = encoded_rhythms.shape[1:]
    n_epochs = 50
    batch_size = 100

    cnn = train_model(
        name='cnn',
        model=CnnModel(input_shape, output_shape),
        x=signals,
        y=encoded_rhythms,
        i_train=i_train,
        i_val=i_val,
        learning_rate=0.0001,
        batch_size=batch_size,
        n_epochs=n_epochs,
    )

    cnn_matrices = evaluate_model(
        model=cnn,
        x=signals,
        y=encoded_rhythms,
        i_train=i_train,
        i_val=i_val,
        i_test=i_test,
        batch_size=batch_size,
    )

```

	Name	Type	Params	Mode	In sizes	Out sizes
0	model	CnnModel	14.2 K	train	[1, 1, 1280]	[1, 5]
1	model.layers	Sequential	14.2 K	train	[1, 1, 1280]	[1, 5]
2	model.layers.0	Conv1d	48	train	[1, 1, 1280]	[1,

8,	1280]						
3	model.layers.1	BatchNorm1d	16	train	[1, 8, 1280]	[1,	
8,	1280]						
4	model.layers.2	ReLU	0	train	[1, 8, 1280]	[1,	
8,	1280]						
5	model.layers.3	MaxPool1d	0	train	[1, 8, 1280]	[1,	
8,	640]						
6	model.layers.4	Conv1d	656	train	[1, 8, 640]	[1,	
16,	640]						
7	model.layers.5	BatchNorm1d	32	train	[1, 16, 640]	[1,	
16,	640]						
8	model.layers.6	ReLU	0	train	[1, 16, 640]	[1,	
16,	640]						
9	model.layers.7	MaxPool1d	0	train	[1, 16, 640]	[1,	
16,	320]						
10	model.layers.8	Conv1d	2.6 K	train	[1, 16, 320]	[1,	
32,	320]						
11	model.layers.9	BatchNorm1d	64	train	[1, 32, 320]	[1,	
32,	320]						
12	model.layers.10	ReLU	0	train	[1, 32, 320]	[1,	
32,	320]						
13	model.layers.11	MaxPool1d	0	train	[1, 32, 320]	[1,	
32,	160]						
14	model.layers.12	Conv1d	10.3 K	train	[1, 32, 160]	[1,	
64,	160]						
15	model.layers.13	BatchNorm1d	128	train	[1, 64, 160]	[1,	
64,	160]						
16	model.layers.14	ReLU	0	train	[1, 64, 160]	[1,	
64,	160]						
17	model.layers.15	AdaptiveAvgPool1d	0	train	[1, 64, 160]	[1,	
64,	1]						
18	model.layers.16	Flatten	0	train	[1, 64, 1]	[1,	
64]							
19	model.layers.17	Linear	325	train	[1, 64]	[1,	
5]							

```

-----
14.2 K    Trainable params
0         Non-trainable params
14.2 K    Total params
0.057     Total estimated model params size (MB)
20        Modules in train mode
0         Modules in eval mode

```

Sanity Checking: | | 0/? [00:00<?, ?it/s]

Training: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

[illegible]

```

Validation: |          | 0/? [00:00<?, ?it/s]
Validation: |          | 0/? [00:00<?, ?it/s]
Validation: |          | 0/? [00:00<?, ?it/s]
Validation: |          | 0/? [00:00<?, ?it/s]
Validation: |          | 0/? [00:00<?, ?it/s]
Validation: |          | 0/? [00:00<?, ?it/s]
Validation: |          | 0/? [00:00<?, ?it/s]
Validation: |          | 0/? [00:00<?, ?it/s]
Validation: |          | 0/? [00:00<?, ?it/s]
Validation: |          | 0/? [00:00<?, ?it/s]
Validation: |          | 0/? [00:00<?, ?it/s]
Validation: |          | 0/? [00:00<?, ?it/s]
Validation: |          | 0/? [00:00<?, ?it/s]
Validation: |          | 0/? [00:00<?, ?it/s]
Validation: |          | 0/? [00:00<?, ?it/s]
Validation: |          | 0/? [00:00<?, ?it/s]
Validation: |          | 0/? [00:00<?, ?it/s]
Validation: |          | 0/? [00:00<?, ?it/s]

```

After the evaluation is finished, we can plot the confusion matrices for the training, validation, and test sets.

```

[10]: def plot_confusion_matrix(c, labels=None, title=None):
        c = np.asarray(c)

        fig = plt.figure(figsize=(5, 4), constrained_layout=True)
        image = plt.imshow(c, cmap='Blues', interpolation='nearest')

        threshold = (c.min() + c.max()) / 2
        for i, j in itertools.product(range(c.shape[0]), repeat=2):
            if c[i, j] < threshold:
                color = image.cmap(image.cmap.N)
            else:
                color = image.cmap(0)
            text = format(c[i, j], '.2g')
            if c.dtype.kind != 'f':
                integer_text = format(c[i, j], 'd')
                if len(integer_text) < len(text):
                    text = integer_text
            plt.text(j, i, text, color=color, ha='center', va='center')

```


```

if labels is not None:
    plt.xticks(np.arange(c.shape[-1]), labels, rotation=45, ha='right')
    plt.yticks(np.arange(c.shape[-1]), labels)
plt.xlabel('Predictions')
plt.ylabel('References')
if title is not None:
    plt.title(title)

def plot_confusion_matrices(matrices, labels):
    for subset in ('train', 'val', 'test'):
        c = matrices[subset]
        accuracy = np.trace(c) / c.sum()
        title = f'{subset.capitalize()} set (accuracy = {accuracy:.3f})'
        plot_confusion_matrix(c, labels=labels, title=title)

plot_confusion_matrices(cnn_matrices, np.unique(rhythms))

```



Val set (accuracy = 0.733)

References	atrial_fibrillation	193	77	5	21	4
	atrial_flutter	122	144	2	28	4
	normal_sinus_rhythm	13	3	219	44	21
	sinus_bradycardia	5	2	15	278	0
	sinus_tachycardia	2	18	12	2	266
		Predictions				
		atrial_fibrillation	atrial_flutter	normal_sinus_rhythm	sinus_bradycardia	sinus_tachycardia

Test set (accuracy = 0.716)

References	atrial_fibrillation	190	84	0	21	5
	atrial_flutter	171	110	2	12	5
	normal_sinus_rhythm	7	4	219	47	23
	sinus_bradycardia	2	0	9	289	0
	sinus_tachycardia	3	16	14	1	266
		Predictions				
		atrial_fibrillation	atrial_flutter	normal_sinus_rhythm	sinus_bradycardia	sinus_tachycardia

## Question 2

Comment the metrics shown in TensorBoard and the confusion matrices. Does the model overfit? Are there some rhythms that are difficult to classify?

**Answer** The model does not overfit (0.71 for test set against 0.76 for validation). Furthermore, by inspecting the confusion matrices, we can observe that there is a cluster with atrial fibrillation and atrial flutter. In fact, these two rhythms are difficult to classify and are often confused with one another in the model.

## Question 3

Define two custom models to classify cardiac rhythms from ECG signals.

You can directly define the layers of the custom models in the following classes.

```
[ ]: class CustomModel1(torch.nn.Module):

    def __init__(self, input_shape, output_shape):
        super().__init__()
        self.input_shape = input_shape
        self.output_shape = output_shape

        kernel_size = 5
        # Implement you own model here.
        # define two twins CNNs
        self.cnn1 = torch.nn.Sequential(
            torch.nn.Conv1d(self.input_shape[0], 8, kernel_size,
padding='same'),
            torch.nn.BatchNorm1d(8),
            torch.nn.ReLU(),
            torch.nn.MaxPool1d(2),

            torch.nn.Conv1d(8, 16, kernel_size, padding='same'),
            torch.nn.BatchNorm1d(16),
            torch.nn.ReLU(),
            torch.nn.MaxPool1d(2),

            torch.nn.Conv1d(16, 32, kernel_size, padding='same'),
            torch.nn.BatchNorm1d(32),
            torch.nn.ReLU(),
            torch.nn.MaxPool1d(2),

            torch.nn.Conv1d(32, 64, kernel_size, padding='same'),
            torch.nn.BatchNorm1d(64),
            torch.nn.ReLU(),
            torch.nn.AdaptiveAvgPool1d(1),
        )
```

```

        self.cnn2 = torch.nn.Sequential(
            torch.nn.Conv1d(self.input_shape[0], 8, kernel_size,
↪padding='same'),
            torch.nn.BatchNorm1d(8),
            torch.nn.ReLU(),
            torch.nn.MaxPool1d(2),

            torch.nn.Conv1d(8, 16, kernel_size, padding='same'),
            torch.nn.BatchNorm1d(16),
            torch.nn.ReLU(),
            torch.nn.MaxPool1d(2),

            torch.nn.Conv1d(16, 32, kernel_size, padding='same'),
            torch.nn.BatchNorm1d(32),
            torch.nn.ReLU(),
            torch.nn.MaxPool1d(2),

            torch.nn.Conv1d(32, 64, kernel_size, padding='same'),
            torch.nn.BatchNorm1d(64),
            torch.nn.ReLU(),
            torch.nn.AdaptiveAvgPool1d(1),
        )

        self.classif = torch.nn.Sequential(
            torch.nn.Flatten(),
            torch.nn.Linear(2*64, self.output_shape[0]),
        )

    def forward(self, x):
        tmp = torch.concat([self.cnn1(x), self.cnn2(x)], dim=1)
        return self.classif(tmp)

class CustomModel2(torch.nn.Module):

    def __init__(self, input_shape, output_shape):
        super().__init__()
        self.input_shape = input_shape
        self.output_shape = output_shape

        # Implement you own model here.
        kernel_size = 3
        self.layers = torch.nn.Sequential(
            torch.nn.Conv1d(self.input_shape[0], 8, kernel_size,
↪padding='same'),
            torch.nn.BatchNorm1d(8),
            torch.nn.ReLU(),

```



```

        torch.nn.MaxPool1d(2),

        torch.nn.Conv1d(8, 16, kernel_size, padding='same'),
        torch.nn.BatchNorm1d(16),
        torch.nn.ReLU(),
        torch.nn.MaxPool1d(2),

        torch.nn.Conv1d(16, 32, kernel_size, padding='same'),
        torch.nn.BatchNorm1d(32),
        torch.nn.ReLU(),
        torch.nn.MaxPool1d(2),

        torch.nn.Conv1d(32, 64, kernel_size, padding='same'),
        torch.nn.BatchNorm1d(64),
        torch.nn.ReLU(),
        torch.nn.AdaptiveAvgPool1d(1),

        torch.nn.Conv1d(64, 32, kernel_size, padding='same'),
        torch.nn.BatchNorm1d(32),
        torch.nn.ReLU(),
        torch.nn.AdaptiveAvgPool1d(1),

        torch.nn.Conv1d(32, 16, kernel_size, padding='same'),
        torch.nn.BatchNorm1d(16),
        torch.nn.ReLU(),
        torch.nn.AdaptiveAvgPool1d(1),

        torch.nn.Flatten(),
        torch.nn.Linear(16, self.output_shape[0]),
    )

    def forward(self, x):
        return self.layers(x)

```

You can train and evaluate the first custom model.

```

[15]: custom1 = train_model(
    name='custom1',
    model=CustomModel1(input_shape, output_shape),
    x=signals,
    y=encoded_rhythms,
    i_train=i_train,
    i_val=i_val,
    learning_rate=0.0001,
    batch_size=batch_size,
    n_epochs=n_epochs,
)

```

```

custom1_matrices = evaluate_model(
    model=custom1,
    x=signals,
    y=encoded_rhythms,
    i_train=i_train,
    i_val=i_val,
    i_test=i_test,
    batch_size=batch_size,
)

plot_confusion_matrices(custom1_matrices, np.unique(rhythms))

```

	Name	Type	Params	Mode	In sizes	Out sizes
-----						
0	model	CustomModel1	28.3 K	train	[1, 1, 1280]	[1, 5]
1	model.cnn1	Sequential	13.8 K	train	[1, 1, 1280]	[1, 64, 1]
2	model.cnn1.0	Conv1d	48	train	[1, 1, 1280]	[1, 8, 1280]
3	model.cnn1.1	BatchNorm1d	16	train	[1, 8, 1280]	[1, 8, 1280]
4	model.cnn1.2	ReLU	0	train	[1, 8, 1280]	[1, 8, 1280]
5	model.cnn1.3	MaxPool1d	0	train	[1, 8, 1280]	[1, 8, 640]
6	model.cnn1.4	Conv1d	656	train	[1, 8, 640]	[1, 16, 640]
7	model.cnn1.5	BatchNorm1d	32	train	[1, 16, 640]	[1, 16, 640]
8	model.cnn1.6	ReLU	0	train	[1, 16, 640]	[1, 16, 640]
9	model.cnn1.7	MaxPool1d	0	train	[1, 16, 640]	[1, 16, 320]
10	model.cnn1.8	Conv1d	2.6 K	train	[1, 16, 320]	[1, 32, 320]
11	model.cnn1.9	BatchNorm1d	64	train	[1, 32, 320]	[1, 32, 320]
12	model.cnn1.10	ReLU	0	train	[1, 32, 320]	[1, 32, 320]
13	model.cnn1.11	MaxPool1d	0	train	[1, 32, 320]	[1, 32, 160]
14	model.cnn1.12	Conv1d	10.3 K	train	[1, 32, 160]	[1, 64, 160]
15	model.cnn1.13	BatchNorm1d	128	train	[1, 64, 160]	[1, 64, 160]

64, 160]						
16   model.cnn1.14	ReLU	0	train	[1, 64, 160]	[1, 64, 160]	
17   model.cnn1.15	AdaptiveAvgPool1d	0	train	[1, 64, 160]	[1, 64, 1]	
18   model.cnn2	Sequential	13.8 K	train	[1, 1, 1280]	[1, 64, 1]	
19   model.cnn2.0	Conv1d	48	train	[1, 1, 1280]	[1, 8, 1280]	
20   model.cnn2.1	BatchNorm1d	16	train	[1, 8, 1280]	[1, 8, 1280]	
21   model.cnn2.2	ReLU	0	train	[1, 8, 1280]	[1, 8, 1280]	
22   model.cnn2.3	MaxPool1d	0	train	[1, 8, 1280]	[1, 8, 640]	
23   model.cnn2.4	Conv1d	656	train	[1, 8, 640]	[1, 16, 640]	
24   model.cnn2.5	BatchNorm1d	32	train	[1, 16, 640]	[1, 16, 640]	
25   model.cnn2.6	ReLU	0	train	[1, 16, 640]	[1, 16, 640]	
26   model.cnn2.7	MaxPool1d	0	train	[1, 16, 640]	[1, 16, 320]	
27   model.cnn2.8	Conv1d	2.6 K	train	[1, 16, 320]	[1, 32, 320]	
28   model.cnn2.9	BatchNorm1d	64	train	[1, 32, 320]	[1, 32, 320]	
29   model.cnn2.10	ReLU	0	train	[1, 32, 320]	[1, 32, 320]	
30   model.cnn2.11	MaxPool1d	0	train	[1, 32, 320]	[1, 32, 160]	
31   model.cnn2.12	Conv1d	10.3 K	train	[1, 32, 160]	[1, 64, 160]	
32   model.cnn2.13	BatchNorm1d	128	train	[1, 64, 160]	[1, 64, 160]	
33   model.cnn2.14	ReLU	0	train	[1, 64, 160]	[1, 64, 160]	
34   model.cnn2.15	AdaptiveAvgPool1d	0	train	[1, 64, 160]	[1, 64, 1]	
35   model.classif	Sequential	645	train	[1, 128, 1]	[1, 5]	
36   model.classif.0	Flatten	0	train	[1, 128, 1]	[1, 128]	
37   model.classif.1	Linear	645	train	[1, 128]	[1, 5]	

-----  
-----  
28.3 K      Trainable params

0	Non-trainable params
28.3 K	Total params
0.113	Total estimated model params size (MB)
38	Modules in train mode
0	Modules in eval mode

```
Sanity Checking: |          | 0/? [00:00<?, ?it/s]
```

```
Training: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |           | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

```
Validation: |          | 0/? [00:00<?, ?it/s]
```

[illegible]

Train set (accuracy = 0.776)

References	atrial_fibrillation	617	247	2	29	5
	atrial_flutter	315	530	2	28	25
	normal_sinus_rhythm	27	10	669	135	59
	sinus_bradycardia	14	6	17	863	0
	sinus_tachycardia	16	41	28	2	813
		atrial_fibrillation	atrial_flutter	normal_sinus_rhythm	sinus_bradycardia	sinus_tachycardia
		Predictions				

Val set (accuracy = 0.733)

References	atrial_fibrillation	194	86	1	17	2
	atrial_flutter	140	134	2	22	2
	normal_sinus_rhythm	15	3	219	39	24
	sinus_bradycardia	5	2	12	281	0
	sinus_tachycardia	3	15	9	1	272
		atrial_fibrillation	atrial_flutter	normal_sinus_rhythm	sinus_bradycardia	sinus_tachycardia
		Predictions				

Test set (accuracy = 0.715)

References	atrial_fibrillation	193	88	1	14	4
	atrial_flutter	185	101	2	11	1
	normal_sinus_rhythm	12	5	209	49	25
	sinus_bradycardia	1	2	4	293	0
	sinus_tachycardia	3	13	7	0	277
		atrial_fibrillation	atrial_flutter	normal_sinus_rhythm	sinus_bradycardia	sinus_tachycardia
		Predictions				

And then the second custom model.

```
[16]: custom2 = train_model(
    name='custom2',
    model=CustomModel2(input_shape, output_shape),
    x=signals,
    y=encoded_rhythms,
    i_train=i_train,
    i_val=i_val,
    learning_rate=0.0001,
    batch_size=batch_size,
    n_epochs=n_epochs,
)

custom2_matrices = evaluate_model(
    model=custom2,
    x=signals,
    y=encoded_rhythms,
    i_train=i_train,
    i_val=i_val,
    i_test=i_test,
    batch_size=batch_size,
```

```
)
```

```
plot_confusion_matrices(custom2_matrices, np.unique(rhythms))
```

	Name	Type	Params	Mode	In sizes	Out sizes
0	model	CustomModel2	16.4 K	train	[1, 1, 1280]	[1, 5]
1	model.layers	Sequential	16.4 K	train	[1, 1, 1280]	[1, 5]
2	model.layers.0	Conv1d	32	train	[1, 1, 1280]	[1, 8, 1280]
3	model.layers.1	BatchNorm1d	16	train	[1, 8, 1280]	[1, 8, 1280]
4	model.layers.2	ReLU	0	train	[1, 8, 1280]	[1, 8, 1280]
5	model.layers.3	MaxPool1d	0	train	[1, 8, 1280]	[1, 8, 640]
6	model.layers.4	Conv1d	400	train	[1, 8, 640]	[1, 16, 640]
7	model.layers.5	BatchNorm1d	32	train	[1, 16, 640]	[1, 16, 640]
8	model.layers.6	ReLU	0	train	[1, 16, 640]	[1, 16, 640]
9	model.layers.7	MaxPool1d	0	train	[1, 16, 640]	[1, 16, 320]
10	model.layers.8	Conv1d	1.6 K	train	[1, 16, 320]	[1, 32, 320]
11	model.layers.9	BatchNorm1d	64	train	[1, 32, 320]	[1, 32, 320]
12	model.layers.10	ReLU	0	train	[1, 32, 320]	[1, 32, 320]
13	model.layers.11	MaxPool1d	0	train	[1, 32, 320]	[1, 32, 160]
14	model.layers.12	Conv1d	6.2 K	train	[1, 32, 160]	[1, 64, 160]
15	model.layers.13	BatchNorm1d	128	train	[1, 64, 160]	[1, 64, 160]
16	model.layers.14	ReLU	0	train	[1, 64, 160]	[1, 64, 160]
17	model.layers.15	AdaptiveAvgPool1d	0	train	[1, 64, 160]	[1, 64, 1]
18	model.layers.16	Conv1d	6.2 K	train	[1, 64, 1]	[1, 32, 1]
19	model.layers.17	BatchNorm1d	64	train	[1, 32, 1]	[1, 32, 1]



20		model.layers.18		ReLU		0		train		[1, 32, 1]		[1, 32, 1]
21		model.layers.19		AdaptiveAvgPool1d		0		train		[1, 32, 1]		[1, 32, 1]
22		model.layers.20		Conv1d		1.6 K		train		[1, 32, 1]		[1, 16, 1]
23		model.layers.21		BatchNorm1d		32		train		[1, 16, 1]		[1, 16, 1]
24		model.layers.22		ReLU		0		train		[1, 16, 1]		[1, 16, 1]
25		model.layers.23		AdaptiveAvgPool1d		0		train		[1, 16, 1]		[1, 16, 1]
26		model.layers.24		Flatten		0		train		[1, 16, 1]		[1, 16]
27		model.layers.25		Linear		85		train		[1, 16]		[1, 5]

-----  
-----

16.4 K	Trainable params
0	Non-trainable params
16.4 K	Total params
0.065	Total estimated model params size (MB)
28	Modules in train mode
0	Modules in eval mode

Sanity Checking: | | 0/? [00:00<?, ?it/s]

Training: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

Validation: | | 0/? [00:00<?, ?it/s]

[illegible]

Validation: | | 0/? [00:00<?, ?it/s]  
 Validation: | | 0/? [00:00<?, ?it/s]  
 Validation: | | 0/? [00:00<?, ?it/s]  
 Validation: | | 0/? [00:00<?, ?it/s]

Train set (accuracy = 0.810)

References	atrial_fibrillation	674	206	3	6	11
	atrial_flutter	340	528	4	9	19
	normal_sinus_rhythm	25	3	753	65	54
	sinus_bradycardia	19	7	35	839	0
	sinus_tachycardia	15	17	16	1	851
		atrial_fibrillation	atrial_flutter	normal_sinus_rhythm	sinus_bradycardia	sinus_tachycardia
		Predictions				

Val set (accuracy = 0.735)

References	atrial_fibrillation	213	65	6	7	9
	atrial_flutter	138	134	2	7	19
	normal_sinus_rhythm	16	0	231	25	28
	sinus_bradycardia	10	6	31	253	0
	sinus_tachycardia	6	11	11	1	271
		atrial_fibrillation	atrial_flutter	normal_sinus_rhythm	sinus_bradycardia	sinus_tachycardia
		Predictions				

Test set (accuracy = 0.726)

References	atrial_fibrillation	216	70	2	7	5
	atrial_flutter	197	89	1	6	7
	normal_sinus_rhythm	7	3	223	31	36
	sinus_bradycardia	6	1	15	278	0
	sinus_tachycardia	4	6	7	0	283
		atrial_fibrillation	atrial_flutter	normal_sinus_rhythm	sinus_bradycardia	sinus_tachycardia
		Predictions				

#### Question 4

How do the two custom models perform? Do they overfit? Do they outperform the first model?

**Answer** Both of our custom models slightly outperform the first model in terms of accuracy, but seem to struggle with `atrial_flutter` and `atrial_fibrillation` signals just like the proposed model.